# The SixDesk Run Environment for SixTrack

E. McIntosh, R. De Maria

Summary

This document replaces the "Run Environment for SixTrack" [1], and describes how massive tracking campaigns can be performed with "SixTrack" [2] starting from a MAD-X input file of the LHC or HL-LHC[1] lattices, the so called mask files. It describes a new set of UNIX BASH or Korn shell scripts which allow the use of the Berkeley Open Infrastructure for Network Computing, BOINC [9]) as an alternative to the Linux LSF batch system. This note is also published and regularly updated on the web page cern.ch/sixtrack-ng/doc/sixdesk/sixdesk_env.html.

Version 27, 18 August 2016.

# 1 Objectives

The principal objective of the SixDesk run environment is to allow a physicist to run a tracking campaign, on either the CERN LSF batch system or BOINC, using the familiar SixTrack run environment on Linux. At the same time, the opportunity was taken to group all user modifiable parameters into two files, sixdeskenv and sysenv, and to speed up run_mad6t, the MAD-X to SixTrack conversion, by running in parallel in batch. The specification of tune scans, fractional amplitude ranges and steps, and various other physics options were also added to the run_six script/command. All SixDesk scripts report activity to the sixdesk.log file and exit with an error code if an unrecoverable error is detected. The messages are optionally sent to STDOUT, your screen, in an interactive session.

# 2 Getting Started

First note that in order to use BOINC you must not only have a CERN account but must also be registered in the AFS protection group boinc:users. This is done with the script add_boinc_user but only E. McIntosh, M. Giovannozzi and IT BOINC administrators are authorised to do so.

---

[1]The HiLumi LHC Design Study is included in the High Luminosity LHC project and is partly funded by the European Commission within the Framework Programme 7 Capacities Specific Programme, Grant Agreement 284404.

All the user modifiable parameters have been collected into the two scripts sixdeskenv, sysenv obviating the need to look through all scripts and make matching changes, simplifying the usage of the scripts, and limiting the risk of error.

It is assumed that there is a large amount (at least 1GB) of disk space available to be used as a workspace which we shall call w1 for illustration. The w1 directory can be a link in $HOME to a directory with the same name, usually on an scratch disk (AFS scratch0, scratch1, etc or in a /scratch directory). Note that by default access to all AFS files will be limited; it is useful to do a fs setacl w1 system:anyuser rl so that everything in an AFS workspace can be read by the support team. In the following it is also assumed that the . directory is in the $PATH environment variable.

A workspace can created by e.g.:

```
cd $HOME
mkdir -p scratch0/w1
ln -s ~/scratch0/w1 w1
cd w1
fs setacl . system:anyuser rl
```

and the SixDesk environment is created by

```
cd $HOME/w1
svn checkout  http://svn.cern.ch/guest/sixdesk/sixjobs
```

or

```
cd $HOME/w1
gtar xvzf /afs/cern.ch/project/sixtrack/src/ \
          SixDesk_run_environment.pro.tgz
```

where the first method will get the latest SVN committed version and the second method uses a link to a tar file which will be updated in parallel with the SVN commits.

This will create a directory subtree starting with the directory sixjobs. An overview of this directory is shown in Appendix A. The latest source and postscript of this document can be found in the directory sixjobs/doc. All the scripts are stored in sixjobs itself and the user should always do a cd w1/sixjobs before executing any commands. It may be necessary to prefix any typed commands with a ./ if your shell PATH does not include the current directory ".". The source version of each script can be found in sixjobs/scripts as will be discussed later.

At this point it is important to check that the LOCALE environment variable LANG has the value C or other English-like setting, since it has been noted that the awk program gives errors if this is not the case. This is normally set, but it it is not the default on some laptops and on MacOS.

# 3   Setting up a Study

To facilitate the description of the various procedures we use the directory w1 as workspace and a study named job_tracking as an illustration. Both names can and should be adapted to the user needs.

A study is controlled by the `sixdeskenv` file that contains all the MAD-X information controlling the simulations: location of the input file (called `.mask` files), the location of the temporary files, the definition of the beam distributions, the number of turns to be tracked or stored and the type of post processing. The corresponding mask file `job_tracking.mask` is included in the release in the `sixjobs/mask` directory. Note that the name of the study is contained in the variable `LHCDescrip` i.e. if we have a study called `lhc_nob1` then we must define it with `export LHCDescrip=lhc_nob1`. It should also be noted that amplitudes and amplitude ranges are specified in beam $\sigma$.

There are several supported types of study, short studies using SixTrack and possibly Sussix, long tracking studies using SixTrack, and DA map production using the SixTrack DA version. Performing a tracking study involves running many jobs for an LHC configuration with many different initial phase space amplitudes and angles and linear tunes. The LHC configuration is defined by a mask file in the `sixjobs/mask` sub-directory and the initial conditions in the `sixdeskenv`, `sysenv` files.

The first step is to edit the `sixdeskenv` script and, if necessary, the `sysenv` script.

NOTA BENE:It is essential to issue a `set_env` command after any modification to these files so that they are saved in the `studies/study` directory as described later. Note also, that as shown below all values must be exported.

The main parameter to be set in the `sixdeskenv` file are:

**export LHCDescrip=job_tracking** the name of the study and of the mask file

**export basedir=/afs/cern.ch/user/$initial/$LOGNAME**

**export scratchdir=/afs/cern.ch/user/$initial/$LOGNAME/scratch0**

**export trackdir=$scratchdir/$workspace** can be re-defined to use an area shared between studies or workspaces;

**export sixtrack_input=$scratchdir/sixtrack_input/$workspace/$LHCDescrip** can be changed to point to, and use, an existing set of `sixtrack_input` files.

The above defaults, apart from the name of the study, are usually satisfactory but, once chosen, cannot be changed easily. Other parameters of the run can be decided later.

At this point:

```
cd ~/w1/sixjobs
set_env
```

will create the full directory structure of the environment. As already mentioned this command should also be used to save modified versions of the `sixdeskenv`, `sysenv` files before doing anything else. The command also reports and logs changes.

While it is recommended to run one study per workspace at the time, most users have multiple studies in the workspace. The command `ls studies` can be used to list them and `set_env "name of the study"` to switch between studies. This switch copies the `sixdeskenv`, `sysenv` files from the `studies` directory to `sixdeskhome` i.e. the current `sixjobs` directory.

An important new feature is that it is possible to execute many commands on a specific study without switching. This facilitates running multiple studies in the same workspace and using commands in batch jobs. Thus, instead of switching, a command can be suffixed with a "study" and optionally a "platform". Examples are shown later. This option is NOT available for the backup or recall, study or workspace, commands described in section 19where a `set_env "study"` must first be performed.

The `print_env` command reports the complete environment to the screen; the command `mywhich` reports a few important values.

Finally note that it is vital that the same versions of SixTrack and MAD-X are used for all cases in a study. Ideally the MAD-X `run_mad6t` jobs should be run on the same type of computer to avoid small numeric differences. Compatible versions of these programs for LSF and BOINC are specified in `sysenv`, the so-called "pro" versions. Any changes will be notified.

In view of some problems uncovered with MAD-X result differences the run_mad6t and `check_mad6t` scripts have been modified to make further checks on the results. See section 5 for details.

# 4 Overview of the Data Structure

The `set_env` will have created the `sixtrack_input`, `studies/study`,`track` and `work` directories, as well as various log-file links and directories. The files `sixtrack_input`, `study`, `track`, `work` are in fact links to the actual directories. An empty file with the name of the study is also created to facilitate SHELL command completion.

The `sixtrack_input` directory will later (after a successful mad6t run) hold all the Six-Track input mother files `fort.3.mother1` and `fort.3.mother2` derived from the mother files in the `control_files` directory as well as a `fort.3.mad` and a `fort.3.aux` and one `fort.2_"seedno".gz`, `fort.8_"seedno".gz`, `fort.16_"seedno".gz` for each seed in the range `istamad` to `iendmad` as defined in `sixdeskenv`. In addition it will hold one `mad.dorun_mad6t*` directory for each run_mad6t command, with one mad6t LSF job `mad6t_"seedno".lsf`, one input file `"study"."seedno"`, one LSF job log `"study"_mad6t_"seedno".log` and one mad6t output file `"study".out."seedno"` for each seed.

The `studies` directory will hold one directory for each study, in turn containing the `sixdeskenv`, `sysenv` files for that study.

The `track` directory will become a hierarchy containing all the SixTrack data and results as follows:

**Level 1 Seed** typically 1 to 60

 **general_input** containing normalised emittance and gamma

**Level 2 simul** for long tracking and/or

 **trans momen** for short/sussix runs

**Level 3 tunex_tuney** e.g. 64.31_59.32

**Level 4 amplitude range** e.g. 18_20 and 18-22 after post-processing

4

**Level 5 turns exponent** e.g. `e5` for $10^5$ turns

**Level 6 the phase space angle** e.g. 67.5

**Level 7** input(links), `fort.2,3,8,16.gz` and the result `fort.10.gz` as well as the LSF jobs and logs

A partial view of the `track` tree is shown in Appendix C. A typical lowest level structure (after the study has been completed), for workspace w1, study job_tracking, long run (simul), seed 1, tunes 64.28_59.31, amplitude range 10_12, $10^5$ turns, phase space angle 1.5, is:

```
~/w1/sixjobs/track/1/simul/64.28_59.31/10_12/e5/1.5
fort.10.gz
fort.16.gz ->
~mcintosh/scratch0/sixtrack_input/w1/job_tracking/fort.16_1.gz
fort.2.gz ->
~mcintosh/scratch0/sixtrack_input/w1/job_tracking/fort.2_1.gz
fort.3.gz
fort.8.gz ->
~mcintosh/scratch0/sixtrack_input/w1/job_tracking/fort.8_1.gz
job_tracking%1%s%64.28_59.31%10_12%5%1.5.log
job_tracking%1%s%64.28_59.31%10_12%5%1.5.lsf
```

Note that input files are in fact links to sixtrack_input, in order to save disk space, apart from the fort.3.gz which is of course different for each job. The result file is fort.10.gz described in section 9. The .lsf and .log files are the LSF job file and the LSF job log file for the particular case.

It should be noted that there is a one to one mapping between the name of a case and the directory where the input and output files are stored. The general form of the name of a case in a `long` run is

```
study%seed%s%tune_range%amplitude_range%turns_exponent%angle
```

In the workspace this case would be found in the directory

```
track/'seed'/simul/'tune_range'/'amplitude_range'/  \
e'turns_exponent'/'angle'
```

To give a specific example, a case named

```
job_tracking%1%s%62.31_60.32%10_12%5%18
```

would be found in the `w1/sixjobs/track` directory for this study in the sub-directory

```
sixjobs/track/1/simul/62.31_60.32/10_12/e5/18
```

this particular case being for seed number 1, tunex and tuney 62.31 and 60.32, amplitude range 10 to 12, $10^5$ turns, angle 18. (The directory simul could also be instead `trans` or `momen` for short/sussix runs and the letter 's' in the name being replaced by 't' or 'm'.)

5

This case name is used in the database, to name the LSF jobs and logs, and is also included in the SixTrack `fort.6` output. Note that as warned in a comment in the `sixdeskenv` file this name must NOT have a % character nor two consecutive underscores.

The `work` directory will contain the database flat files for reporting and managing all the tasks and jobs of a particular study. The most interesting files are `work/taskids`, `work/completed_cases`, `work/incomplete_cases` and the subdirectory `work/lsfjobs` (or `boincjobs`). The master file is called `taskids` and contains one line for every case of a study with the case name followed by one or more LSF or BOINC taskids. The `run_status` command reports on the status of the study as described in section 8.

There are also the other directories:

**bin** containing links to various utility programs used by `run_six` (`dalie4 dalie6 readda`), `run_join10` (`joinf10`) and `run_post` (`read10b`)

**control_files** the SixTrack mother files for the collision/injection/beam 2 cases which are used to generate the SixTrack input files

**doc** The Latex source and Postscript of this document

**inc** other mask files and the prepare_fort.3 script

**mask** the LHC description mask files for MAD-X

**plot** the plotting mask files and scripts in the sub-directory `plot/inc`. The sub-directory `plot/"study"` will contain a hierarchy matching the `track` hierarchy, but with the `angle` directory replaced by a hidden directory `.angle` containing some or all of the files `averem.eps.gz distance.eps.gz fort.10.gz fort.15.gz fort.30.gz maxslope.eps.gz smear.eps.gz survival.eps.gz test.ps.gz`.

**sixdeskTaskIds** the study TaskIds for BOINC

**scripts** the source and copies of the commands

**utilities** the various LSF job masks for `run_mad6t` and `run_six`

# 5 Running mad6t to produce the basic SixTrack input files

NOTA BENE: It is essential to have a $LHCDescrip.mask file in the subdirectory `mask` in order to run MAD-X to produce the SixTrack input files. This mask file in turn references many LHC Database files and this often requires some checking. Sample mask files can be found in the directory mask:

```
jobref503_withbb_coll.mask    jobrefslhc_withbb_coll.mask
jobref503_withbb_inj.mask     jobref503_inj.mask
jobrefslhc_inj.mask           jobrefslhc_withbb_inj.mask
jobref503_coll.mask           jobrefslhc_coll.mask
job_tracking.mask
```

The file job_tracking.mask is copied from /afs/cern.ch/eng/lhc/optics/SLHCV3.01 and all the others from /afs/cern.ch/eng/lhc/optics/SLHCV2.0 [10].

In order to investigate a series of random seeds the particular seed number in the MAD input file has to be replaced by a variable name:

```
Set, SEEDSYS , 1 ;
Set, SEEDRAN , 1 ;
```

becomes

```
Set, SEEDSYS , %SEEDSYS ;
Set, SEEDRAN , %SEEDRAN ;
```

The place keepers `%SEEDSYS` and `%SEEDRAN` will be replaced automatically by a proper seed number by `run_mad6t` based on the current seed number based on the `istamad` and `iendmad` variables defined in `sixdeskenv`.

In addition, in order to use the correct bunch size, the place keeper `%NPART` is replaced by the value specified for the `bunch_charge` variable also in `sixdeskenv`.

Please note that the MAD-X `sixtrack` command takes its information from the last `Twiss, save` command. It is a sensible precaution to put these commands consecutively in the MAD script and mask files.

The version of MAD-X to be used for the conversion runs is defined in `sysenv` and defaults to the current production version. At this point it is also necessary to specify in `sixdeskenv`

**pmass=938.272046** The mass of the proton [4] which can be reset to 938.271998 for backwards compatibility with earlier studies.

**bunch_charge=1.1500e+11** New bunch_charge variable for fort.3.mother1_[col/inj]

**runtype=** inj or coll for injection/collision for an LHC lattice. For more information on how to define runtypes for either the LHC or other machines please see section 20.

**beam=** null or b1 or B1 for Beam1, b2 or B2 for Beam2.

**CORR_TEST=** 0 or 1 if check_mad6t is to copy the corrector strengths for each seed into one file in sixtrack_input.

**fort_34=** If null the fort.34 files will not be copied to the `sixtrack_input` directory. These files define the multipole strengths for the linear lattice when doing Second Order Detuning and Distortion (SODD) analysis in expert mode.

**istamad=1** first seed for MAD-X.

**iendmad=** normally 60 (maximum 64) but it is recommended to use iendmad=1 until the results of the `run_mad6t` are considered satisfactory. The seeds are used to generate variations of the basic LHC machine with different randomly generated magnet errors.

**madclass=** `8nm` for 8 normalised CPU minutes, or `1nh` for 1 normalised hour for the `run_mad6t` LSF jobs where 8nm is often enough but `1nh` may be necessary.

After a `set_env`, a `run_mad6t` or `run_mad6t -i` may be performed. The `-i` option means run interactively and get the output to the screen or redirected to a file and can be useful for testing the mask file and mad6t run. The `platform` option is ignored as mad6t runs are either performed on the desktop or on LSF. For reasons of numerical compatibility all the `run_mad6t` runs should be performed on the same type of machine. Subsequently, LSF jobs may be used by run_mad6t, one per seed, and they will run in parallel. In the case of multiple studies in a workspace, either a `set_env "study"` can be performed or the study can be specified on the command as `run_mad6t` "study". In all cases, the success/failure/correctness of these runs should be verified. The script `check_mad6t` checks what it can but it is essential to have a look at the MAD-X output. This output can be found in the most recent `sixtrack_input/mad.run_mad6t*` directory. Once the MAD-X run has completed successfully for one seed, `iendmad` can be set to the desired value (typically 60/64) and a new `run_mad6t` command performed. Every `run_six`, see later, starts with an internal `check_mad6t`. To facilitate checking the mad6t1.lsf and mad6t.lsf scripts now do some additional checking. Two files ERRORS and WARNINGS may be produced and can be found in sixtrack_input. Ideally, after initial tests, the contents of sixtrack_input should be deleted and one `run_mad6t` should be done for all seeds to ensure consistency. `run_six` will not run if the ERRORS file exists.

# 6   Problems, cleaning up, and support

There are often problems at this initial stage: at any time the command `rm -r work/*` `track/* sixtrack_input/*` will completely clean up and allow a restart from the beginning. Help and diagnostics and error messages will be found in the `sixdesk.log` file, one per study [1].

# 7   run_six - Launching SixTrack Runs

It is usual to commence a study with some short runs or a few long runs using LSF, perhaps with just a few seeds.

The script `run_six` is used after setting the various variables as explained below. It will automatically launch tracking jobs into the LSF batch system using the batch scripts in the directory `utilities` or into the BOINC job submission buffer. It is recommended to use LSF batch for short runs and for exploratory studies, or when the binary files are required for detailed examination. BOINC is recommended for production runs and for large studies of more than a few thousand jobs. Note that BOINC does not produce the fort.20 graphics nor return any binary files. Only the `fort.10.gz` result file is returned which is enough for the subsequent post-processing. In addition note that the platforms LSF and BOINC should not both be used in a single study; rather, once the preliminary investigations are complete using LSF, a new study can be cloned to run full scale production with BOINC using the same `sixtrack_input files` thus avoiding another redundant `run_mad6t`. The experienced user is free to modify the script as explained

---

[1]Help with these procedures is always available from Eric.McIntosh@cern.ch, by telephone, or by SKYPE to mcintosh94. Accelerator Physics issues should be initially addressed to Massimo.Giovannozzi@cern.ch or Riccardo.de.Maria@cern.ch.

later. The following variables, all exported, as specified in the `sixdeskenv` file are used by this and subsequent scripts.

**tunex & tuney** The required horizontal and vertical tunes. These may be different for collision and injection as shown later.

**emit** The normalised LHC emittance, used for the computation of the beam $\sigma$.

**emit_beam** The normalised LHC beam emittance for substitution into the mask (replaces the keyword `%EMIT_BEAM`) and into the mother file (replaces the keyword `%emit_beam`).

**xing** Half crossing angle in $\mu$rad for substitution into the mask (replaces the keyword `%XING`) and into the mother file (replaces the keyword `%xing`).

**e0** Energy of reference particle depending on runtype

**dpini & dpmax** At injection the initial relative momentum deviation is set to 'dpini=0.00075' and at top energy it is set to 'dpini=0.00027'. For the determination of the (non-linear) chromaticity a wider range is used: 'dpmax=0.002' (see below).

**kstep** Used to define the step width of the phase space angle. For further information see section 7.1. The phase space angle is related to the emittance ratio via $\phi = \arctan\left(\sqrt{\epsilon_y/\epsilon_x}\right)$, where the emittance is defined as $\epsilon_z = A_z^2/\beta_z$, for $z = x, y$.

**dimen** The dimensionality of phase space can be chosen between 4 and 6. In the latter case the full six–dimensional tracking is done including cavities.

**chrom** To correct for slight differences between MAD and SixTrack the chromaticity is routinely corrected by setting 'chrom' to 1 and using

**chrom_eps=0.0000001** This operation will not be performed for 'chrom=0' but `chromx` `chromy` will be used instead.

**chromx=2.** and

**chromy=2.** being the values used when 'chrom=0'.

**sussix** To determine precise values for the detuning calculation this switch should be set to: 'sussix=1'. It uses the sussix program [5]. This option is only valid for the short run configuration (see section 7.1). (Problems have been found when using `sussix` in the 6D case.)

    `run_six` can handle three different modes of tracking: Normally initial investigations are carried out with `short` runs (and possibly `sussix`).

1. Short run — This run mode is used to find chromaticity and detuning as a function of $\delta$ and amplitude respectively. Typically this is done with just 1,000 turns (activate with `short=1` in `sixdeskenv`). The other variables in `sixdeskenv` for this run are described in section 7.1.

2. Long run — This mode is meant for the dynamic aperture determination proper (activate with `long=1` in `sixdeskenv`). The other variables for this run are described in section 7.2.

3. Differential Algebra (DA) run — If high order Taylor maps are needed this is the mode to use (activate with `da=1`). This mode is mostly for expert use. The `SixTrack` author F. Schmidt should be consulted on how to make best use of it. In this case `run_six` calls the porgrams readda, dalie4 and dalie6 in the directory `bin`.

Note that only one type of run may be chosen at any one time; one and only one of the `sixdeskenv` variables `short, long, da` may be set to 1.

The short/long runs both use seeds as specified by

**ista=$istamad** Start seed

**iend=$iendmad** End seed

As shown above, the default seed range `ista, iend` for `run_six` is the same as that used for `run_mad6t`, namely `istamad, iendmad`. These values may be changed at any time, for example to submit jobs for a limited range of seeds, but must clearly be a subset of the `run_mad6t` values.

## 7.1 Short Run

**ns1s & ns2s** Lower and upper amplitude range in beam $\sigma$.

**nss** Amplitude step in beam $\sigma$.

**turnss** Number of turns which is usually set to 1,000 in this mode.

**turnsse** This variable should be set to the number of zeros of 'turnss', i.e. '3' in our example, it becomes part of the data directory structure. Therefore, if one decides to redo this analysis at say 10,000 turns one specifies 'turnsse=4' and subsequently the data are stored separate from those produced with 'turnsse=3'.

**writebins** This defines after how many turns data are written to output files. In this mode it should always be set to: 'writebins=1' since all turns are needed to find the tunes as a function of amplitude.

**kini & kend** Initial and end angle in phase space. Typically set from '1' to 'kmax=5' (see next variable). By specifying 'kini=0' the nonlinear chromaticity is calculated as well (which uses the 'dpmax' setting) and thereafter the initial angle is set back to: 'kini=1'. Note that the variation from 'kini' to 'kend' is done in steps defined by 'kstep'.

**kmax** This defines the number of phase space angles, e.g. 'kmax=5' means that each step `kstep` is of: $90°/(kmax + 1) = 15°$.

**reson=0** switch for Guignard resonance calculation

## 7.2   Long Run

**ns1l & ns2l**  Lower and upper amplitude range in beam $\sigma$. This range is sub–divided into ranges of `nsincl` $\sigma$. In each job 30 pairs of particles are evenly distributed in each subrange. The close–by pairs are used to find the onset of chaos. Typically we find that a variation 2 $\sigma$ is sufficiently dense to find the minimum dynamic aperture with a precision of 0.5 $\sigma$.

**nsincl**  2 $\sigma$ is standard. A smaller step, of 0.5 say, can give better results.

**turnsl**  For the long term tracking we usually track for 100,000 turns or for one miliion.

**turnsle**  This variable should be set to the number of zeros of 'turnsl', i.e. '5' in our example.

**writebinl**  This defines after how many turns data are written to output files.

   *Important*: make sure that `writebinl` is large enough otherwise huge amounts of data will be created. Occasionally that may be of use, however in most cases make sure that no more than a total of 1,000 turns are recorded. This implies that for `turnsl=100000` the variable should be set to at least `writebinl=100`'. When running on BOINC, rather than LSF, the binary files are not returned but SixTrack is checkpointed every `writebinl` turns. It is recommended to set `writebinl` to 10000 for BOINC runs of 100,000 turns or more.

**kinil & kendl**  Initial and end angle in phase space. As in the 'short run' mode the variation from 'kinil' to 'kendl' is done in steps defined by 'kstep'.

**kmaxl**  This defines the number of phase space angles, e.g. 'kmaxl=5' means that each steps amounts to: $90°/(kmaxl+1) = 15°$. Thus the actual angles are computed by dividing 90 by kmaxl+1, so 5, 19, 59 for example are reasonable choices. The choice of `kmaxl` is discussed in Ref. [6].

   Now the other physics and system parameters must be defined in the `sixdeskenv` file if the defaults are not suitable. When the platform is defined as LSF, LSF job class definitions will be required:

**platform=LSF**  or may be set to BOINC.

**longlsfq=1nd**  sufficient for 100,000 turns, 60 particles

**classs=sixmedium**  for short runs.

**classda=sixda**  for the sixda jobs requiring large memory

**sixdeskforce=0**  Should normally be left at 0 but may be set to 1 or 2 (see later)

The LSF job class defaults are normally satisfactory, but `longlsfq` should be set to 2nd or to 1nw if performing more than $10^5$ turns.

**ibtype=0**  or 1 to use the Erskine/McIntosh optimised error function of a complex number

**idfor=1**  the closed orbit is added, if set to 0 the initial co-ordinates are unchanged

**sixdeskpairs=30**  Normal value for 60 particles, naximum of 32.

Then we have, depending on the `runtyp`:

```
if test $runtype = "inj"
then
```

**e0=450000.**  (energy)

**gamma=479.6**  (gamma)

**dpini=0.00075**  (initial relative momentum deviation)

```
elif test $runtype = "col"
then
```

**e0=7000000.**

**gamma=7460.5**

**dpini=0.00027**

```
fi
```

**dpmax=0.002**  maximum momentum deviation for a short term run

Next we have the tunes again depending on the runtype:

**tune=0**

In this case the `run_six` will make a special local LSF run to compute the tunes. Alternatively the tunes may be specified and in particular a tunescan can be performed where the tunes will be computed on a straight line from (tunex,tuney) with gradient deltay/deltax up to and including (tunex1,tuney1). The tunes must be 10 .le. tune .lt. 100 in format dd.dd[d][d]. The folowing example specifies the tunes (64.28,59.31) with no scan (because `tunex = tunex1, tuney = tuney1`.

```
if test $runtype = "inj"
then
```

**tunex=64.28**  Start value

**tuney=59.31**  Start value

**deltax=0.001**  Increment to tunex

**deltay=0.001**  Increment to tuney

**tunex1=64.28**  End value

**tuney1=59.31**  End value

12

```
elif test $runtype = "col"
then
```

**tunex=64.31**

**tuney=59.32**

**deltax=0.001**

**deltay=0.001**

**tunex1=64.31**

**tuney1=59.32**

```
fi
```

The total number of jobs/cases can be computed as the product of the number of seeds, the number of tune pairs, the number of amplitude intervals and the number of angles. The total number of cases and progress is reported in the work directory and can be examined with the `run_query` or `run_status` commands.

Each batch job returns a result file `fort.10.gz` to the `track` tree/hierarchy. For LSF studies, if the CASTOR switch is on in the `sysenv` file, all the result files, including the binary files and the `fort.6` output, are compressed in a tar file and written to `$CASTOR_HOME/direct_track/*` where the `direct_track` tree matches the `track` tree in AFS.

It is not recommended to run a study of more than 30,000 jobs/cases in a single workspace but up to 100,000 still works. A study can be split over two or more workspaces, perhaps by seed number, and the results combined.

Problems may, indeed often, arise when it is necessary to run the script more than once, either because of a system crash, batch daemon not responding, jobs lost, or some other error. The script `run_six` can be rerun as often as necessary, but by default will not re-submit jobs. The `run_six` can either be submitted as a batch job and/or with a different seed range (`ista, iend`). The `run_six` script also keeps a copy of each LSF job in the `track` hierarchy along with the data files or links `fort.2.gz fort.3.gz fort.8.gz fort.16.gz`. It also updates the `taskids` file in the `sixjobs/work` directory. The `taskids` file contains one line for each case. Each line contains the JobName (which can be mapped to the case directory) and the associated LSF job_ID(s) or BOINC ID(s). If `run_six` is called more than once, it ignores cases where a non-zero fort.10.gz exists, but otherwise deletes the tracking input files and re-generates them. Even simpler recovery is available with the `run_missing_jobs` command. *It is essential to wait for existing LSF jobs to complete, or to cancel them, and do at least two `run_status` commands before invoking `run_missing_jobs`.* This procedure has proven to be extremely effective. In the worst of all cases wait for, or cancel, outstanding batch jobs with the `bkill 0` command, and delete everything except the `sixtrack_input` with a `rm track/* work/*` command and restart.

With BOINC the situation is simpler; near the end of a study a `run_incomplete_cases` command will re-submit jobs for those cases where no results are available. The script `run_incomplete_cases_lsf` can also be used to submit LSF jobs instead to get rid of the annoying tail of unfinished Work Units.

# 8  Monitoring the progress of the study.

The script `run_status` looks into the database `work` directory and updates and summarises the status of the study. The `run_query` command gives a quick look without any updating. It first counts the number of cases in `work/taskids`. It reports the number of LSF batch jobs generated, possibly more than one per case, and does an LSF `bjobs` to report on job status. Finally it performs a rather time-consuming search to find the number of unfinished cases. It also produces the files `completed_jobs`, `incomplete_jobs` and possibly `missing_jobs` in the `$sixdeskwork` directory. When all cases are complete the run_join10 procedure can be initiated. While the `run_status` command may take some time, many minutes for 50,000 cases, it makes it very easy indeed to recover from LSF failures. For example, in a recent study with 43,100 jobs over 800 jobs failed. One `run_missing_jobs` command re-submitted them automatically to successfully complete the tracking. As an alternative, the `run_six` command can be re-issued and will not re-submit jobs if `sixdeskforce = 0`. It will not re-submit jobs which have been completed successfully if `sixdeskforce = 1` AND a `fort.10.gz` result file has been created.

There are also cases of database corruption. These are cleaned up, after all running LSF jobs have terminated, by the `correct_cases` command. Overall the environment has been proven to be rather robust and almost all errors can be recovered as long as there is a valid `work/taskids` file.

# 9  The `fort.10` file

The structure of the `fort.10` files is shown in Tables 1 and 2, which have been taken from the official SixTrack manual [2] and updated with recent additions. All values are double-precision numbers, values encoded if necessary. There is one (very long) line per particle.

# 10  The Other SixTrack Output Files and What the Run Environment Does With Them.

By default the scripts automatically tar all output files and store them directly to CASTOR. In order to view these files they have to be copied from CASTOR using the standard CASTOR tools like `rfcp`, `rfrm` and `nsls, nsrm, nsmkdir` etc. The option `--help` is available for each of these commands.

Although the full description of all output files is beyond the scope of this note; an ascii file worth looking at is the `fort.6.gz` which gives an explicit description of all operations and possible failures of the SixTrack run as seen by the program. Also one finds the tracking data in the binary files `fort.90.gz fort.89.gz etc` down to `fort.59.gz` in the cases of 64 particles/32 pairs, which may be useful for further analysis. These and other files are all described in the SixTrack manual [2]. Note that these files are not available from BOINC runs which are for full scale production; LSF should be used for testing and debugging.

Table 1: Post-processing data of the `fort.10` file

| # of Column | Description |
|---|---|
| 1 | Maximum turn number |
| 2 | Stability Flag (0=stable, 1=lost) |
| 3 | Horizontal Tune |
| 4 | Vertical Tune |
| 5 | Horizontal $\beta$–function |
| 6 | Vertical $\beta$–function |
| 7 | Horizontal amplitude $1^{st}$ particle |
| 8 | Vertical amplitude $1^{st}$ particle |
| 9 | Relative momentum deviation $\frac{\Delta p}{p_o}$ |
| 10 | Final distance in phase space |
| 11 | Maximum slope of distance in phase space |
| 12 | Horizontal detuning |
| 13 | Spread of horizontal detuning |
| 14 | Vertical detuning |
| 15 | Spread of vertical detuning |
| 16 | Horizontal factor to nearest resonance |
| 17 | Vertical factor to nearest resonance |
| 18 | Order of nearest resonance |
| 19 | Horizontal smear |
| 20 | Vertical smear |
| 21 | Transverse smear |
| 22 | Survived turns $1^{st}$ particle |
| 23 | Survived turns $2^{nd}$ particle |
| 24 | Starting seed for random generator |
| 25 | Synchrotron tune |
| 26 | Horizontal amplitude $2^{nd}$ particle |
| 27 | Vertical amplitude $2^{nd}$ particle |

Table 2: Post-processing data of the `fort.10` file continued

| # of Column | Description |
|:---:|:---:|
| 28 | Minimum horizontal amplitude |
| 29 | Mean horizontal amplitude |
| 30 | Maximum horizontal amplitude |
| 31 | Minimum vertical amplitude |
| 32 | Mean vertical amplitude |
| 33 | Maximum vertical amplitude |
| 34 | Minimum horizontal amplitude (linear decoupled) |
| 35 | Mean horizontal amplitude (linear decoupled) |
| 36 | Maximum horizontal amplitude (linear decoupled) |
| 37 | Minimum vertical amplitude (linear decoupled) |
| 38 | Mean vertical amplitude (linear decoupled) |
| 39 | Maximum vertical amplitude (linear decoupled) |
| 40 | Minimum horizontal amplitude (nonlinear decoupled) |
| 41 | Mean horizontal amplitude (nonlinear decoupled) |
| 42 | Maximum horizontal amplitude (nonlinear decoupled) |
| 43 | Minimum vertical amplitude (nonlinear decoupled) |
| 44 | Mean vertical amplitude (nonlinear decoupled) |
| 45 | Maximum vertical amplitude (nonlinear decoupled) |
| 46 | Emittance Mode I |
| 47 | Emittance Mode II |
| 48 | Secondary horizontal $\beta$–function |
| 49 | Secondary vertical $\beta$–function |
| 50 | $Q'_x$ |
| 51 | $Q'_y$ |
| 52 | SixTrack Version (encoded in double precision) |
| 53 – 58 | Closed Orbit |
| 59 | The number of the Random Set |
| 60 | Tracking CPU time in seconds |

# 11 Post-processing with run_join10

run_join10 gathers the results of completed jobs and produces combined output files at level 4 in the track tree. It also deletes any results of previous run_join10 commands. The result directory is named $ns1l-$ns2l e.g. 14-20 as compared to the other amplitude directories e.g. 14_16, 16_18, etc. There is one additional parameter for run_join10 namely turnsemax which is set automatically by sixdeskenv to $turnsle or $turnsse for long/short studies. Missing files are reported to sixdesk.log but even incomplete results may be useful. This procedure makes use of the utility program joinf10 in the directory bin.

# 12 Post-processing with run_post

After possibly joining the fort.10 files they can be postprocessed to find chaotic boundaries and particle losses by using the script run_post, which in turn uses the program read10b in the directory bin.

There are a couple of options for plotting. The sixdeskenv variables iplot and kvar specific to run_post are used. If iplot=1 a plot is produced for each seed and the results can be found in the plot directory. Note that setting iplot to 1 will produce a possibly huge amount of data, even if it is compressed with gzip. The variable kvar should be set to 1 (the default) to obtain the DA as a function of angles for a long study, and the DA over ALL seeds and angles is plotted for each angle even if iplot=0. Note that this is an extremely time consuming procedure. (For special post-processing of the fort.10.gz files the set of fort.15 files are produced and saved when iplot = 1)

The following variables are used, almost certainly with the same values as used by run_six and run_join10:

**kinil & kendl** Initial and end angle in phase space. The variation from 'kinil' to 'kendl' is done in steps defined by 'kstep' (see below).

**kmaxl** This defines the number of phase space angles, e.g. 'kmaxl=5' means that each step is of: $90°/(kmaxl + 1) = 15°$.

**kstep** Used to define the step width of the phase space angle.

**Ampl** The amplitude range in sigma. To distinguish the original amplitudes ranges from the joined ranges a hyphen is used '–' instead of an underscore '_'. Ampl is set automatically using ns1s/ns2s or ns1l/ns2l.

**turnse** This variable should be set to the number of zeros of number of turns processed, i.e. '5' in our example, it is part of the data directory structure.

**short** In the example 'short=0' means that the mode short run is not activated.

**long** In the example 'long=1' means that the mode long run is activated.

**iplot** No plotting for each seed if 'iplot=0'. If this flag is set to '1' or '2' the following graphics are produced:

- Short run

  - Chromaticity, i.e. the tune versus $\delta$.
  - Horizontal and vertical detuning each in one plot for all tracked phase space angles.
  - Tune foot print, i.e. vertical tune versus horizontal tune with the amplitude as a parameter.

- Long run

  - End value of the distance in phase space 'd(turns)' of 2 initially close–by particles as a function of initial amplitude.
  - Fitted slope of $\log d(turns)$ versus $\log(turns)$ of the distance in phase space of 2 initially close–by particles as a function of initial amplitude. For details of the meaning of these two chaotic definitions please refer to reference [8].
  - Survival plot, i.e. survival time versus initial amplitude.
  - Horizontal and vertical smear as a function of initial amplitude.
  - Phase space averaged amplitude versus initial amplitude.

For 'iplot=2' these plots are automatically printed using your normal Linux $PRINT_CMD. Obviously, great care has to be taken to avoid a swamping the printer. The graphics are stored as files like `test.ps.gz` in the directory `sixjobs/plot/"study"` with the same tree structure as the `track` tree but with a hidden `.angle` directory.

The `DA*` result files, one per angle, from `run_post` in long mode are called `DAres."Study"."Tunes."Turns exponent"."Angle"`, and are stored in the `sixjobs` directory itself. The contents and format are defined in the next section 13 It should be noted again that setting `kvar=1` causes `run_post` to take a long time, even a very long time of many hours, for a large number of angles. If both `iplot` and `kvar` are set to 0 the `DAres.Study.Tunes.Turns exponent.Angle` files are produced anyway and can be processed with `run_awk` or your own procedures. A partial view of the `sixjobs/plot` directory is shown in Appendix C.

# 13   The DAres files

These files contain the following columns:

- Run name for particular seed.

- "Strict" chaotic boundary via slope method [8].

- "Certain" chaotic boundary via large distance in phase space method [8].

- Dynamic aperture concerning the phase space averaged amplitude (preferred value).

- Raw dynamic aperture concerning initial amplitude (to be used with care).

- Lower bound of tracked amplitude range.

- Upper bound of tracked amplitude range.

The algorithm producing the values can be described in the following steps. The code (`read10b`) first reads the joined `fort.10`. If the emittance of the mode I (`emitI`) is smaller than `1E-10` then the initial coordinate $x$ is replaced by `sqrt(betx*emitI)`. The same for $y$.

The first amplitude of the `fort.10` is analyzed in order to obtain:

- `rat`: the tangent of the `emitII/emitI`;

- `rad`: ratio between `emitI/emit` ( or `emitII/emit` if `rat>1`) if
  `emit= emitI**2+emitII**2` normalized with the real emittance of the beam specified by the user. If `emitI` is small then `emitI` is estimated using the initial conditions;

- `rad1`: same quantity as `rad` but evaluated as before using the average amplitudes if `emitI` is small otherwise using the average amplitudes and the secondary beta functions;

`rat`, `rad`, and `rad1` will be used for all the other amplitudes in the following analysis. The analysis of all the amplitudes starts with:

- the chaotic boundary with slope method, `achaos`,

- distance in the phase space with slope method, `achaos1`

- initial condition of lost particle, `alost2`

- average initial condition, average corrected, of a lost particle, `alost1`

`achaos` is set to the first amplitude (normalized with `rad`) until an amplitude is found chaotic. It means that if the first amplitude is chaotic or all amplitudes are chaotic the result does not change. `achaos1` is zero until an amplitude is found chaotic. `alost2` is the amplitude normalized with `rad` of the first lost amplitude.

For the amplitudes between the `achaos` and `achaos1`, the average ratio between `rad*x` and `rad1*x_average` is computed. If the ratio is in between 0.9 and 1.1 than `alost2` is corrected by this average ratio and stored in `alost1`. If outside of the bounds `alost1` is inverted. If there is no amplitudes in between `achaos` and `achaos1`, `alost1==alost2`.

The code ends reporting `achaos`, `achaos2`, `alost1`, `alost2`, `rad` normalized first amplitude and last amplitude.

`run_awk` uses the DAres* files from `run_post`, reports to the screen, and produces a simple plot file named DAres*.plot.

# 14 Running studies in batch and/or parallel

In summary a study can be run in the following steps:

Create a study by `set_env` and after creating the `mask` file in the `mask` directory do a `run_mad6t`. Repeat as often as necessary until `check_mad6t` is successful and the `mad6t` output has been verified.

For a large study use `run_six` in a batch job by for example:

```
bsub -q1nd 'cd w1/sixjobs;run_six job_tracking LSF'
```

to submit all cases of the study `job_tracking` in workspace `w1` using LSF. The progress of the command can be monitored by `tail -f sixdesk.log` or `bpeek "LSFJobID"`.

1. Use `run_status` or `run_status job_tracking`, as frequently as convenient until all cases are complete or all batch jobs have terminated. The command `correct_cases` can be used if the database appears to be inconsistent (but only after all LSF jobs have terminated). It checks each case and updates the complete/incomplete status. If there are incomplete cases with LSF do a `run_missing_jobs`. With BOINC use a `run_incomplete_cases` to speed up the completion of the study or to handle the tail of incomplete cases. When all cases have been completed do a

2. `run_join10` or `bsub -q1nd 'cd w1/sixjobs;run_join10 job_tracking'` followed by a

3. `run_post` or `bsub -q1nd 'cd w1/sixjobs;run_post job_tracking'`

# 15  Lockfiles

It is necessary to use locks in order to avoid conflicting modifications to a file. A directory is locked if it contains a file sixdesklock in read only mode (444). If a script fails or dies for some reason or a batch job is killed this lock file may be left behind. The script `check_locks` reports lock status for the current study; `check_all_locks` for the workspace. If the study itself is locked, it will be reported first as such a lock inhibits further checking. `unlock "directory"` frees the lock. The `unlock_all` unlocks all locks, but is deprecated when there is more than one active study in the workspace. If in doubt about a lock, and the information shown by the check commands, simply do a `check_lock "directory"` or `cat sixdesklock` and `ls -l sixdesklock` in the relevant directory, to see which script, process and machine are holding the lock and since when. The general philosophy is to lock the study so that only one operation at a time is permitted. Most command wait for the study to be unlocked but the `run_results` command for BOINC just exits on the grounds that it will be run periodically, nornally as an `acrontab` entry.

# 16  BOINC

The sixdesk environment was designed to make it as transparent as possible to use LSF or the Berkeley Open Infrastructure for Network Computing BOINC [9]. which has made over 100,000 PCs available for tracking studies. There is one additional step required, namely `run_results`, to retrieve result files from the BOINC Web server and move the `fort.10.gz` to the appropriate directory. Only this file is returned, no `fort.6` or binary or graphics files and nothing is written to CASTOR. To use BOINC the `sixdeskenv` file is modified to

**platform=BOINC**

20

and that is all.

1. `run_mad6t` operates identically as for LSF.

2. `check_mad6t` as for LSF.

3. `run_six` as for LSF.

4. `run_status` as for LSF.

5. `run_results` must now be called to get the fort.10.gz file from the BOINC server. It can be called regularly, automatically, by using an acrontab entry in AFS, and there is an example `acrontab.entry` in the `sixjobs` directory. In addition a `run_incomplete_cases` script has been implemented to allow the use of multiple tasks for a case. This is particularly useful towards the end of a run, when say 90% of the cases are complete, in order to speed up completion of the study.

   Unlike LSF which returns all result files directly to the `track` tree with no checking and over-writing existing results, BOINC compares the results file fort.10.gz with any existing result and reports any differences to `sixdesk.log`.

6. when all cases are complete the run_join10 and `run_post` procedures are used as with LSF.

# 17    Modifying the scripts/commands

All modifications should be made to the scripts in the directory `scripts`. The source script beginning with the characters "my" should be modified when it exists rather than the derived script e.g. modify `myrun_six` and NOT run_six. The script `domyseds`, which operates on all the files/scripts in the file `allscripts`, should then be executed to expand macros in scripts named "my*" to produce the actual command scripts, and to copy all the commands to `../sixjobs`.

Here is a more or less alphabetic list of the "my" scripts:

- mybackup_study to CASTOR

- mybackup_workspace to CASTOR

- mycheck_all_locks Tries to check for any locks in the workspace

- mycheck_lock Checks a specified directory or `sixjobs`

- mycheck_locks Tries to check for any locks in the current study

- mycheck_mad6t Does minimal checks on the `sixtrack_input`

- mycorrect_cases Re-generates completed_cases and `incomplete_cases` using the `taskids` file and checking for the presence of `fort.10.gz`

- mydelete_study deletes all or parts of a study based on replies to prompts

- mydorun_mad6t called by run_mad6t

21

- myget_all_betavalues Writes all betavalues to STDOUT

- mymywhich Writes some important variables to STDOUT

- myrecall_study from CASTOR

- myrecall_workspace from CASTOR

- myrerun_all_cases (Used to rerun_all cases on BOINC)

- myrun_awk Processes all the DAres* files and writes a summary to STDOUT

- myrun_incomplete_cases Re-submits incomplete cases to BOINC

- myrun_incomplete_tasks (Rarely used)

- myrun_join10 Joins all `fort.10.gz` files in preparation for `run_post`

- myrun_missing_jobs Re-submits any LSF jobs which have failed (after `run_status`

- myrun_post Does the post-processing to produce graphics and the DAres* files

- myrun_query A quick and brief status report

- myrun_results Retrieves BOINC results

- myrun_six Submits all LSF jobs or BOINC Work Units (after run_mad6t

- myrun_status Reports status and updates the database for the platform LSF

- myset_env Used to save the environment after changes to `sixdeskenv sysenv`, ito create a study, or to switch studies

- myunlock Unlocks the specified directory or `sixjobs`

- myunlock_all An unlock of all locks of the current study and any other workspace locks not specific to a study `sixjobs, plot` and `sixdeskTaskIds`

and here is a list of some other useful scripts/commands without macros.

- bresume_all Resume all suspended batch jobs

- bstop_all Stop all batch jobs

- get_wpro A sample script for using from acron to get results

- minav.awk The awk script used by run_awk

- print_env Writes the `sixdeskenv sysenv` variables and values to STDOUT

- query_all Does a `run_status` for each study in a workspace

- sub_wpro A sample script for doing a `run_six` in batch LSF

# 18 The Subroutines

These subroutines are all defined in the file mydot_profile and may therefore use macros themselves. A `. ./dot_profile` statement is issued when starting a script in order to make them available. They are called by the corresponding "my" macros as shown.

- sixdeskmess() mymess `messagelevel text` where the messagelevel is an integer 0, 1 or 2 and the text is a doublequoted string. It uses the `sixdeskenv` parameters `sixdeskecho` and `sixdesklevel`. If `sixdesklevel` = 0 then only basic and error messages are logged. If `sixdesklevel` = 1 then additional information is logged and if `sixdesklevel` = 2 then further debug ifnormation is also logged. If `sixdeskecho` = "" only minimum information is echoed to STDOUT and otherwise all messages are also echoed.

- sixdeskmktmp() mymktmp `name (dir)` makes a uniquely named temporary file in the specified directory `dir` or in `.` and returns the full pathname in `$name`

- sixdeskmktmpdir() mymktmpdir `name (dir)` makes a uniquely named temporary directory in the specified directory `dir` or in `.` and returns the full pathname in `$name`

- sixdeskexit() myexit `integer` exits with the code `integer` freeing any locks

- sixdeskunlock() myunlock `(dir)` unlocks the specified directory or `.`

- sixdesklock() mylock `(dir)` Locks the specified directory `(dir)` or `.`

- sixdesktunes() mytunes Converts the tunes and deltas specified in `sixdeskenv` to integers and logs the values in `sixdesk.log`

- sixdeskinttunes() myinttunes Uses the integer tunes to produce values to replace placeholders in fort.3

- sixdeskamps() myamps Generates the equivalent integer values for amplitudes

- sixdeskrundir() myrundir `name1 name2` Converts the specified run `name1` to the equivalent unique directory name which is returned in `$name2`

The scripts:

  my  dot_boinc

  my  dot_bsub

are effectively subroutines used by run_six. The scripts

  my  dot_env

  my  dot_profile

are called by almost every other script to establish the environment and make the subroutines available.

# 19   Backing up the workspace

The AFS scratch disks at CERN are not (yet) backed up. It is recommended that you periodically run a backup so that data can be recovered easily in case of a scratch disk failure. It is also recommended to a `backup_study [job_tracking]` when a study `job_tracking` is complete. The study may then be deleted safely.

These backups can be done rather simply in a batch job. First is shown an example of backing up your scratch disks and then the rather more sophisticated facilities available for SixDesk tracking studies and workspaces.

NOTA BENE: all `castor_backups` save the logs in your `/castor/logs`. Although these are gzipped and might be useful if you forget the contents of a backup you may also just want to delete them to save space in your HOME directory. The contents of backups can be always be determined from the backups themselves. Note also that for all backups, links are NOT followed. Further a link to a non-existent file is NOT backed up.

## 19.1   Backing up scratch disks

Everyone has a CASTOR account and a default `$CASTOR_HOME` directory `/castor/cern.ch/user/$INITIAL/$LOGNAME`. All the `ns` commands like `nsls`, `nsrm`, `nsmkdir` etc use this by default. In this example all the backups are in the CASTOR directory `scratch_backups`. In the examples below `date` is of the form `yyyy-mm-dd` (and `mydate` is just `-yyyy-mm-dd`). The batch job script, which must have execute permission, is as follows:

```
#!/bin/sh
# Do a dated (by day) backup of all scratch disks to
#  $CASTOR_HOME/scratch_backups
# and log the reports in ~/backup_scratch.reports
#
cd $HOME
mydate="-`date -Idate`"
echo "" >> backup_scratch.reports
echo "Scratch backup reports$mydate" >> backup_scratch.reports
echo "Scratch backup reports$mydate"
echo "" >> backup_scratch.reports
for i in `ls -d scratch*`
do
  echo "Backing up $i to scratch_backups/$i$mydate" \
       >> backup_scratch.reports
  echo "Backing up $i to scratch_backups/$i$mydate"
  castor_backup $i scratch_backups/$i$mydate
  nsls -l scratch_backups/$i$mydate >> backup_scratch.reports
  nsls -l scratch_backups/$i$mydate
done
```

If the job runs to completion (as can be checked in the LSF STDOUT or the backup_scratch_reports) a lost disk, scratch0 say, can be restored with a

```
cd $HOME
castor_recall scratch_backups/scratch0"mydate"/scratch0
```

where "mydate" identifies the backup to be used. `nsls scratch_backups` will show all available backups. Clearly this can also be done in a batch job. To recall to a different place

```
cd $HOME/scratch99
castor_recall scratch_backups/scratch0"date"/scratch0 .
```

will recall all the files to the directory `scratch99`.

## 19.2   Backing up SixDesk workspaces and directories

The intended usage is that in order to free up disk space a study can be backed up and then deleted. It may later be recalled to exactly the same workspace or to a new different workspace. The `backup_workspace` command might be used to backup a complete set of studies. A study may be recalled from either a workspace backup or a study backup.

Deleted studies are never backed up; a non-deleted study cannot be recalled to the same workspace unless it has first been deleted. Deleting a study does NOT delete the `sixjobs` directory and all `sixdeskenv` and `sysenv` files are kept in the studies directory. All DAres* files are also preserved.

All backup/recall commands should be issued from the `sixjobs` directory as usual except for `recall_workspace` which should be one level above e.g. `w1`.

In these examples "date" is of the form ddmmyy. So, normally, for example:

```
cd ˜w1/sixjobs
set_env job_tracking
backup_study
delete_study
```

and subsequently

```
cd ˜w1/sixjobs
set_env job_tracking
recall_study w1%job_tracking%"date"
```

or to recall to a new different workspace

```
cd ˜w1/sixjobs
set_env job_tracking
recall_study w1%job_tracking%"date" w99
```

In the latter case the sixdeskenv files are edited to reflect the new workspace. In all recalls, an existing sixjobs directory is never overwritten, but in all cases the sixtrack_input, track, work, and plot data are recalled along with logfiles if possible. To recall several studies to a new workspace, the first recall will recall sixjobs as well as the data; afterwards, `cd` to the workspace to recall other studies. As mentioned a study can be similarly recalled from a workspace backup by

```
      set_env job_tracking
      recall_study w1%"date" to restore it to w1
```

or

```
      set_env job_tracking
      recall_study w1%"date" w99 to a new empty workspace.
```

Finally a complete workspace can be recalled to the same workspace or a different one. To backup a workspace

```
      cd ~/w1/sixjobs
      backup_workspace
```

and to recall

```
      cd ~/w1
      recall_workspace w1%"date"
```

or, since a new different workspace is empty, we must first get a copy of the recall_workspace script

```
      cd ~/w99
      cp SOMEWHERE/recall_workspace .
      ./recall_workspace w1%"date"
```

All backups can be found in $CASTOR_HOME/workspace_backups with names like w1%job_tracking%"date" for a study backup or w1%"date" for a workspace backup. The command nsls workspace_backups will list all available backups.

All backups are automatically restartable from checkpoints as they may take some time and are subject to the usual system failures. A file backup_study.list or backup_workspace.list is used for restarting from the point of failure. Thus a backup must be completed before a new backup is started in the same workspace (or the file backup_study.list or backup_workspace.list must be deleted). The workspace is locked to ensure only one backup at a time and no switching of studies.

These procedures may seem complicated but the simple backup, delete, recall are easy to perform and the complications are necessary to avoid destroying data and to recover from system failures. Details of the CASTOR backups/ recalls can be found in the castor_backup.log or castor_recall.log in the sixjobs directory.

# 20  Other lattices

Clearly any lattice file may be processed by the run environment not just the LHC. The mask file has to be changed as appropriate. Other changes are described in the main part of this note, including aesthetic name changes. Also required are definitions of the runtype (see section 5). This variable and the variable beam specify which of the fort.3.mother1 and fort.3.mother2 files from the directory sixjobs/control_files to use. The value of the runtype and beam variables simply specify the suffix of the files to be used from this directory. For how these files are constructed the user is refered to the SixTrack manual [2].
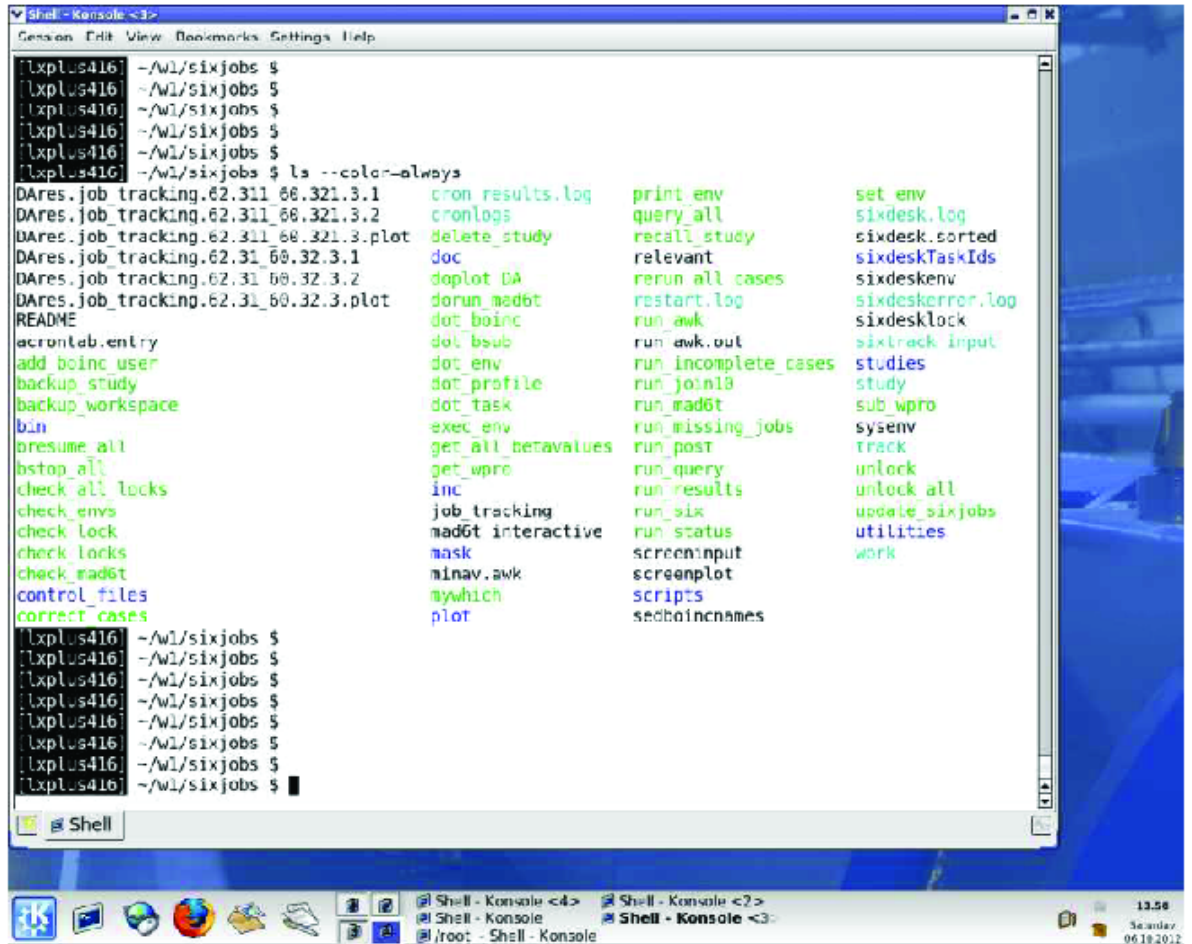
26

Figure 1: Example of fu problems.

# 21 Acknowledgements

Thanks to F. Schmidt for the original scripts and many helpful discussions and much advice; M. Giovannozzi and T. Risselada for acting as guinea pigs and for many helpful suggestions. Finally, thanks to S. Fartoukh for setting up the `mask` files in the `mask` directory.

# A An overview of the w1/sixjobs directory

See Figure 1

# B    A partial view of the track tree

For the purpose of illustration a very small study with just two seeds, two tune pairs, two amplitude and two angles is used.

```
Level 1
track: 1 2 general_input

Level 2
track/1: simul

Level 3
track/1/simul: 62.311_60.321 62.31_60.32

Level 4
track/1/simul/62.31_60.32: 6-7 6.5_7 6_6.5 \
betavalues mychrom sixdesktunes

Level 5
track/1/simul/62.31_60.32/6-7: e3

Level 6
track/1/simul/62.31_60.32/6-7/e3: 30 60

Level 7
track/1/simul/62.31_60.32/6-7/e3/30:
fort.10.gz
track/1/simul/62.31_60.32/6-7/e3/60:
fort.10.gz

Level 6
track/1/simul/62.31_60.32/6_6.5/e3: 30 60

Level 7
track/1/simul/62.31_60.32/6_6.5/e3/30:
fort.10.gz
fort.16.gz
fort.2.gz
fort.3.gz
fort.8.gz
job_tracking%1%s%62.31_60.32%6_6.5%3%30.log
job_tracking%1%s%62.31_60.32%6_6.5%3%30.lsf

track/1/simul/62.31_60.32/6_6.5/e3/60:
fort.10.gz
fort.16.gz
```

```
fort.2.gz
fort.3.gz
fort.8.gz
job_tracking%1%s%62.31_60.32%6_6.5%3%60.log
job_tracking%1%s%62.31_60.32%6_6.5%3%60.lsf
```

Note the correspondance between the LSF jobname and the name of the directory containing it, as described in Section 3.

## C  A partial view of the plot tree after post-processing

In this view post-processing was performed with `iplot=1` and `kvar=1`. Again, a very small study with just two seeds, two tune pairs, two amplitudes and two angles is used.

```
plot: inc job_tracking
job_tracking: 1 2 allseeds.eps.gz

job_tracking/1: simul

job_tracking/1/simul: 62.311_60.321 62.31_60.32

job_tracking/1/simul/62.31_60.32: 6-7

job_tracking/1/simul/62.31_60.32/6-7/e3: .1 .2

job_tracking/1/simul/62.31_60.32/6-7/e3/.1:
averem.eps.gz
distance.eps.gz
fort.10.gz
fort.11
fort.15.gz
fort.28
fort.30.gz
fort.40
kvar.eps.gz
maxslope.eps.gz
smear.eps.gz
survival.eps.gz
test.ps.gz

job_tracking/1/simul/62.31_60.32/6-7/e3/.2:
averem.eps.gz
distance.eps.gz
fort.10.gz
fort.11
```

```
fort.15.gz
fort.28
fort.30.gz
fort.40
kvar.eps.gz
maxslope.eps.gz
smear.eps.gz
survival.eps.gz
test.ps.gz
```

Note the "hidden" directories `.1` and `.2` as mentioned in Section .

# D   Alphabetic list of supported commands and scripts

- backup_study to CASTOR

- backup_workspace to CASTOR

- bresume_all Resume all suspended batch jobs

- bstop_all Stop all batch jobs

- check_all_locks Tries to check for any locks in the workspace

- check_lock Checks a specified directory or `sixjobs`

- check_locks Tries to check for any locks in the current study

- check_mad6t Does minimal checks on the `sixtrack_input`

- correct_cases Re-generates completed_cases and `incomplete_cases`

- delete_study deletes all or parts of a study based on replies to prompts

- dorun_mad6t called by run_mad6t

- get_all_betavalues Writes all betavalues to STDOUT

- get_wpro A sample script for using from acron to get results

- mywhich Writes some important variables to STDOUT

- print_env Writes the `sixdeskenv sysenv` variables and values to STDOUT

- query_all Does a `run_status` for each study in a workspace

- recall_study from CASTOR

- recall_workspace from CASTOR

- rerun_all_cases (Used to rerun_all cases on BOINC)

- run_awk Processes all the DAres* files and writes a summary to STDOUT

- run_incomplete_cases Re-submits incomplete cases to BOINC

- run_incomplete_tasks (Rarely used)

- run_join10 Joins all `fort.10.gz` files in preparation for `run_post`

- run_missing_jobs Re-submits any LSF jobs which have failed (after `run_status`

- run_post Does the post-processing to produce graphics and the DAres* files

- run_query A quick and brief status report

- run_results Retrieves BOINC results

- run_six Submits all LSF jobs or BOINC Work Units (after run_mad6t

- run_status Reports status and updates the database for the platform LSF

- set_env Used to save the environment after changes to `sixdeskenv sysenv`, to create a study, or to switch studies

- sub_wpro A sample script for doing a `run_six` in batch LSF

- unlock Unlocks the specified directory or `sixjobs`

- unlock_all An unlock of all locks of the current study and any other workspace locks not specific to a study `sixjobs, plot` and `sixdeskTaskIds`

# References

[1] M. Hayes and F. Schmidt, "Run Environment for SixTrack", LHC Project Note 300. http://cern.ch/frs/SixTrack_run_environment/manual.ps.

[2] F. Schmidt, "SixTrack: Version 3, Single Particle Tracking Code Treating Transverse Motion with Synchrotron Oscillations in a Symplectic Manner, User's Reference Manua", CERN/SL/94–56 (AP). cern.ch/frs/Documentation/doc.html

[3] H. Grote, "Statistical significance of dynamic aperture calculations", Beam Physics Note 34.

[4] http://physics.nist.gov/cuu/Constants/index.html.

[5] R. Bartolini and F. Schmidt, "SUSSIX: A Computer Code for Frequency Analysis of Non–Linear Betatron Motion", presented at the workshop "Nonlinear and Stochastic Beam Dynamics in Accelerators – A Challenge to Theoretical and Computational Physics", Lüneburg, September 29 – October 3, 1997, CERN SL/Note 98–017 (AP), http://cern.ch/frs/report/sussix_manual_sl.pdf.

[6] M. Giovannozzi and E. McIntosh, "Parameter scans and accuracy estimates of the dynamical aperture of the CERN LHC", EPAC'06, June 2006, Edinburgh.

[7] M. Böge and F. Schmidt, "Data Organisation for the LHC Tracking Studies with SIX-TRACK", LHC Project Note 99,
http://cern.ch/frs/report/lhc_pro_note99.pdf.

[8] M. Böge and F. Schmidt, "Estimates for Long–Term Stability for the LHC", LHC Project Report 114, presented in part at the Particle Accelerator Conference, Vancouver, 12–16 May, (1997), AIP Conference Proceedings 405 (1996), http://cern.ch/frs/report/conj97lhc.pdf, the poster version: http://cern.ch/frs/report/conj97post.pdf
and contribution to the workshop on "New Ideas for Particle Accelerators", Santa Barbara, November 1996.

[9] "Berkeley Open Infrastructure for Network Computing", http://boinc.berkeley.edu

[10] S. Fartoukh, "Optics and layout cern repository for the LHC upgrade": /afs/cern.ch/eng/lhc/optics/SLHCV2.0/, /afs/cern.ch/eng/lhc/optics/SLHCV3.01/ .